



**Министерство науки и высшего образования Российской Федерации
Федеральное государственное бюджетное образовательное учреждение
высшего образования
«Московский государственный технический университет
имени Н.Э. Баумана
(национальный исследовательский университет)»
(МГТУ им. Н.Э. Баумана)**

ФАКУЛЬТЕТ _____ «Информатика и системы управления»

КАФЕДРА _____ «Теоретическая информатика и компьютерные технологии»

Лабораторная работа № 5
по курсу «Теория искусственных нейронных сетей»
«Сверточные нейронные сети (CNN)»

Студентка группы ИУ9-72Б Самохвалова П. С.

Преподаватель Каганов Ю. Т.

Москва 2023

1 Задание

1. Реализовать нейронную сеть LeNet. В качестве базы данных использовать MNIST. Применить оптимизаторы SGD, AdaDelta, NAG, Adam.
2. Реализовать нейронную сеть VGG16. В качестве базы данных использовать CIFAR-10. Применить оптимизаторы SGD, AdaDelta, NAG, Adam.
3. Реализовать нейронную сеть ResNet (34). В качестве базы данных использовать ImageNet. Применить оптимизаторы SGD, AdaDelta, NAG, Adam.

2 Практическая реализация

Исходный код программы представлен в листингах 1- 3.

Листинг 1: LeNet

```
1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import torchvision
5 import torchvision.transforms as transforms
6 from torch.utils.data import DataLoader
7 import torch.nn.functional as F
8 import matplotlib.pyplot as plt
9
10
11 transform = transforms.Compose([
12     transforms.ToTensor(),
13     transforms.Normalize((0.1307,), (0.3081,))
14 ])
15
16
17 trainset = torchvision.datasets.MNIST(root='./data', train=True,
18     download=True, transform=transform)
19
20 testset = torchvision.datasets.MNIST(root='./data', train=False,
21     download=True, transform=transform)
22
23
24 trainloader = DataLoader(trainset, batch_size=64, shuffle=True)
25 testloader = DataLoader(testset, batch_size=64, shuffle=False)
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
```

```

26         super(LeNet, self).__init__()
27         self.conv1 = nn.Conv2d(1, 6, kernel_size=5)
28         self.conv2 = nn.Conv2d(6, 16, kernel_size=5)
29         self.fc1 = nn.Linear(16*4*4, 120)
30         self.fc2 = nn.Linear(120, 84)
31         self.fc3 = nn.Linear(84, 10)
32
33     def forward(self, x):
34         x = F.relu(self.conv1(x))
35         x = F.max_pool2d(x, 2)
36         x = F.relu(self.conv2(x))
37         x = F.max_pool2d(x, 2)
38         x = x.view(-1, 16*4*4)
39         x = F.relu(self.fc1(x))
40         x = F.relu(self.fc2(x))
41         x = self.fc3(x)
42         return x
43
44
45 model = LeNet()
46
47 criterion = nn.CrossEntropyLoss()
48
49 optimizer_sgd = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
50 optimizer_adadelta = optim.Adadelta(model.parameters())
51 optimizer_nag = optim.SGD(model.parameters(), lr=0.01, momentum=0.9,
52                             nesterov=True)
52 optimizer_adam = optim.Adam(model.parameters())
53
54
55 def train_test_model(optimizer, name):
56     model.train()
57     losses = []
58     for epoch in range(10):
59         running_loss = 0.0
60         for i, data in enumerate(trainloader, 0):
61             inputs, labels = data
62             optimizer.zero_grad()
63             outputs = model(inputs)
64             loss = criterion(outputs, labels)
65             loss.backward()
66             optimizer.step()
67             running_loss += loss.item()
68         epoch_loss = running_loss / len(trainloader)
69         losses.append(epoch_loss)
70         print(f"{name} - Epoch {epoch + 1} loss: {epoch_loss}")

```

```

71
72     model.eval()
73     correct = 0
74     total = 0
75     with torch.no_grad():
76         for data in testloader:
77             images, labels = data
78             outputs = model(images)
79             _, predicted = torch.max(outputs, 1)
80             total += labels.size(0)
81             correct += (predicted == labels).sum().item()
82     accuracy = 100 * correct / total
83     print(f"{name} - Accuracy: {accuracy}%")
84
85     return losses
86
87
88 sgd_losses = train_test_model(optimizer_sgd, "SGD")
89 adadelat_losses = train_test_model(optimizer_adadelat, "AdaDelta")
90 nag_losses = train_test_model(optimizer_nag, "NAG")
91 adam_losses = train_test_model(optimizer_adam, "Adam")
92
93
94 epochs = range(1, 11)
95 plt.plot(epochs, sgd_losses, label='SGD')
96 plt.plot(epochs, adadelat_losses, label='AdaDelta')
97 plt.plot(epochs, nag_losses, label='NAG')
98 plt.plot(epochs, adam_losses, label='Adam')
99 plt.xlabel('Epoch')
100 plt.ylabel('Loss')
101 plt.title('Loss function dependence on the number of epochs for each
           optimizer')
102 plt.legend()
103 plt.show()

```

Листинг 2: VGG16

```

1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import torchvision
5 import torchvision.transforms as transforms
6 import matplotlib.pyplot as plt
7
8 model = torchvision.models.vgg16(pretrained=False)
9
10 num_classes = 10

```

```

11 model.classifier[6] = nn.Linear(4096, num_classes)
12
13 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
14 model.to(device)
15
16 transform = transforms.Compose(
17     [
18         transforms.ToTensor(),
19         transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
20     ]
21 )
22
23 trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
24     download=True, transform=transform)
25
26 trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
27     shuffle=True, num_workers=2)
28
29 testset = torchvision.datasets.CIFAR10(root='./data', train=False,
30     download=True, transform=transform)
31
32 testloader = torch.utils.data.DataLoader(testset, batch_size=4, shuffle=
33     False, num_workers=2)
34
35 optimizers = {
36     "SGD": optim.SGD(model.parameters(), lr=0.001, momentum=0.9),
37     "AdaDelta": optim.Adadelta(model.parameters(), lr=0.01),
38     "NAG": optim.SGD(model.parameters(), lr=0.001, momentum=0.9,
39         nesterov=True),
40     "Adam": optim.Adam(model.parameters(), lr=0.00001)
41 }
42
43 for optimizer_name, optimizer in optimizers.items():
44     criterion = nn.CrossEntropyLoss()
45     epochs = 5
46
47     optimizer_losses = [0] * epochs
48
49     for epoch in range(epochs):
50         running_loss = 0
51         for i, data in enumerate(trainloader, 0):
52             inputs, labels = data
53             inputs, labels = inputs.to(device), labels.to(device)
54
55             optimizer.zero_grad()
56
57             outputs = model(inputs)
58             loss = criterion(outputs, labels)

```

```

52         loss.backward()
53         optimizer.step()
54
55         running_loss += loss.item()
56
57         optimizer_losses[epoch] += loss.item()
58
59         if i % 2000 == 1999:
60             print(f"[{optimizer_name}, {epoch + 1}, {i + 1}] loss: {
running_loss / 2000}")
61             running_loss = 0.0
62
63         optimizer_losses[epoch] /= len(trainloader)
64
65         correct = 0
66         total = 0
67         with torch.no_grad():
68             for data in testloader:
69                 images, labels = data
70                 images, labels = images.to(device), labels.to(device)
71                 outputs = model(images)
72                 _, predicted = torch.max(outputs, 1)
73                 total += labels.size(0)
74                 correct += (predicted == labels).sum().item()
75
76         print(f"Accuracy of the network with {optimizer_name} optimizer:
{100 * correct / total}%")
77
78         plt.plot(range(1, epochs + 1), optimizer_losses, label=
optimizer_name)
79         plt.xlabel('Epoch')
80         plt.ylabel('Loss')
81         plt.title('Loss function dependence on the number of epochs for
optimizer ' + optimizer_name)
82         plt.legend()
83         plt.show()

```

Листинг 3: ResNet (34)

```

1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import torchvision.transforms as transforms
5 import torchvision.datasets as datasets
6 import torchvision.models as models
7 import matplotlib.pyplot as plt
8

```

```

9
10 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
11
12
13 transform = transforms.Compose([
14     transforms.Resize(256),
15     transforms.CenterCrop(224),
16     transforms.ToTensor(),
17     transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
18         0.225])
19 ])
20
21 train_dataset = datasets.CIFAR10(root='./data', train=True, download=
22     True, transform=transform)
23 test_dataset = datasets.CIFAR10(root='./data', train=False, download=
24     True, transform=transform)
25
26 train_loader = torch.utils.data.DataLoader(dataset=train_dataset,
27     batch_size=64, shuffle=True)
28 test_loader = torch.utils.data.DataLoader(dataset=test_dataset,
29     batch_size=64, shuffle=False)
30
31 model = models.resnet34(pretrained=False)
32 model.to(device)
33
34 criterion = nn.CrossEntropyLoss()
35
36 accuracies = {}
37
38
39 optimizers = ['SGD', 'Adadelata', 'NAG', 'Adam']
40 losses = {optimizer_name: [] for optimizer_name in optimizers}
41
42 for optimizer_name in optimizers:
43     if optimizer_name == 'SGD':
44         optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
45     elif optimizer_name == 'Adadelata':
46         optimizer = optim.Adadelata(model.parameters())
47     elif optimizer_name == 'NAG':
48         optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9,
49             nesterov=True)

```

```

49     elif optimizer_name == 'Adam':
50         optimizer = optim.Adam(model.parameters(), lr=0.001)
51
52     for epoch in range(5):
53         model.train()
54         running_loss = 0.0
55         for images, labels in train_loader:
56             images, labels = images.to(device), labels.to(device)
57             optimizer.zero_grad()
58             outputs = model(images)
59             loss = criterion(outputs, labels)
60             loss.backward()
61             optimizer.step()
62             running_loss += loss.item()
63         print(f"Epoch {epoch+1}, Optimizer: {optimizer_name}, Loss: {
running_loss / len(train_loader)}")
64         losses[optimizer_name].append(running_loss / len(train_loader))
65
66     model.eval()
67     correct = 0
68     total = 0
69     with torch.no_grad():
70         for images, labels in test_loader:
71             images, labels = images.to(device), labels.to(device)
72             outputs = model(images)
73             _, predicted = torch.max(outputs, 1)
74             total += labels.size(0)
75             correct += (predicted == labels).sum().item()
76     accuracy = 100 * correct / total
77     accuracies[optimizer_name] = accuracy
78     print(f'Accuracy of the network on the test images with {
optimizer_name} optimizer: {accuracy:.2f}%')
79
80     plt.plot(range(1, 6), losses[optimizer_name], label=optimizer_name)
81     plt.xlabel('Epoch')
82     plt.ylabel('Loss')
83     plt.title('Loss function dependence on the number of epochs for
optimizer ' + optimizer_name)
84     plt.legend()
85     plt.show()
86
87
88 print("Accuracies for different optimizers:")
89 for optimizer_name, accuracy in accuracies.items():
90     print(f"{optimizer_name}: {accuracy:.2f}%")

```


3 Результаты

Результаты работы программы представлены на рисунках 1- 19.

```
SGD - Epoch 1 loss: 0.04797668762797396
SGD - Epoch 2 loss: 0.036027234870811865
SGD - Epoch 3 loss: 0.028856778872211893
SGD - Epoch 4 loss: 0.0244847006027939
SGD - Epoch 5 loss: 0.021115803868930996
SGD - Epoch 6 loss: 0.01876699934248874
SGD - Epoch 7 loss: 0.014785314569834805
SGD - Epoch 8 loss: 0.014005225215688177
SGD - Epoch 9 loss: 0.010436913913792856
SGD - Epoch 10 loss: 0.010134605349389519
SGD - Accuracy: 98.86%
AdaDelta - Epoch 1 loss: 0.022558863759092957
AdaDelta - Epoch 2 loss: 0.019058083317939987
AdaDelta - Epoch 3 loss: 0.016210800011888387
AdaDelta - Epoch 4 loss: 0.013096760662922034
AdaDelta - Epoch 5 loss: 0.012577806868266005
AdaDelta - Epoch 6 loss: 0.010397543301558959
AdaDelta - Epoch 7 loss: 0.009481603296815095
AdaDelta - Epoch 8 loss: 0.00808082507300579
AdaDelta - Epoch 9 loss: 0.006468135072998011
AdaDelta - Epoch 10 loss: 0.005467295557532801
AdaDelta - Accuracy: 98.93%
```

Рис. 1 — Результат работы нейронной сети LeNet с оптимизаторами SGD и AdaDelta

```
NAG - Epoch 1 loss: 0.00787582927497022
NAG - Epoch 2 loss: 0.004608031559518695
NAG - Epoch 3 loss: 0.0022396754083366184
NAG - Epoch 4 loss: 0.003136958142601099
NAG - Epoch 5 loss: 0.003943243743814708
NAG - Epoch 6 loss: 0.0016330818259314116
NAG - Epoch 7 loss: 0.00037742435931199456
NAG - Epoch 8 loss: 0.0006183253060741904
NAG - Epoch 9 loss: 0.00020005558085133064
NAG - Epoch 10 loss: 4.119652412146766e-05
NAG - Accuracy: 99.16%
Adam - Epoch 1 loss: 0.01374092939565212
Adam - Epoch 2 loss: 0.011992538895982336
Adam - Epoch 3 loss: 0.009622250434174022
Adam - Epoch 4 loss: 0.008653295581118986
Adam - Epoch 5 loss: 0.009547748116000702
Adam - Epoch 6 loss: 0.0103140324600208
Adam - Epoch 7 loss: 0.007459128912773552
Adam - Epoch 8 loss: 0.008228338507701181
Adam - Epoch 9 loss: 0.006582019857332815
Adam - Epoch 10 loss: 0.006534842348275204
Adam - Accuracy: 98.93%
```

Рис. 2 — Результат работы нейронной сети LeNet с оптимизаторами NAG и Adam

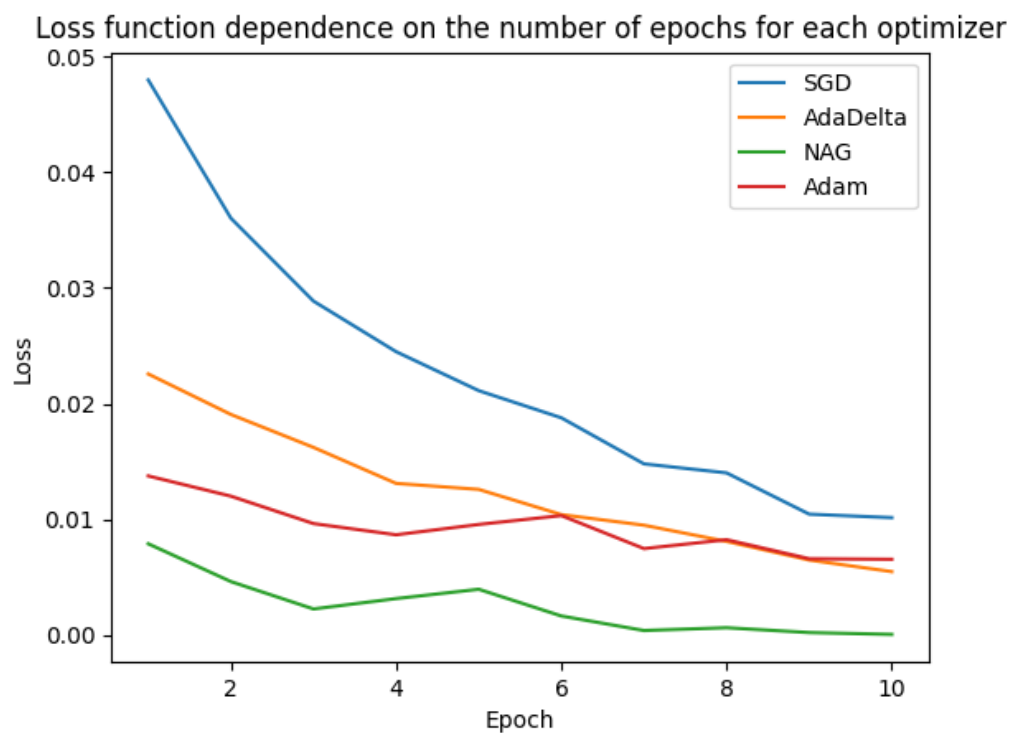


Рис. 3 — Результат работы нейронной сети LeNet с оптимизаторами SGD, AdaDelta, NAG и Adam




Рис. 4 — Результат работы нейронной сети VGG16 с оптимизатором SGD

```

[SGD, 1, 2000] loss: 2.1894213616847993
[SGD, 1, 4000] loss: 1.9679649308621883
[SGD, 1, 6000] loss: 1.8337315746247769
[SGD, 1, 8000] loss: 1.7300100834965706
[SGD, 1, 10000] loss: 1.6710455544143916
[SGD, 1, 12000] loss: 1.5788842151761056
[SGD, 2, 2000] loss: 1.4742421067953109
[SGD, 2, 4000] loss: 1.407913557395339
[SGD, 2, 6000] loss: 1.3453061136119067
[SGD, 2, 8000] loss: 1.2827470286041498
[SGD, 2, 10000] loss: 1.22350149115175
[SGD, 2, 12000] loss: 1.149056606795173
[SGD, 3, 2000] loss: 1.0750381919629872
[SGD, 3, 4000] loss: 1.0461009511547164
[SGD, 3, 6000] loss: 0.9965575462942943
[SGD, 3, 8000] loss: 0.9744769980029087
[SGD, 3, 10000] loss: 0.950291696835775
[SGD, 3, 12000] loss: 0.929938881078735
[SGD, 4, 2000] loss: 0.811417927344286
[SGD, 4, 4000] loss: 0.8250379110233044
[SGD, 4, 6000] loss: 0.8186864303138282
[SGD, 4, 8000] loss: 0.783252324091387
[SGD, 4, 10000] loss: 0.7893385257571935
[SGD, 4, 12000] loss: 0.7868428526644129
[SGD, 5, 2000] loss: 0.6670214962890968
[SGD, 5, 4000] loss: 0.6642861536218697
[SGD, 5, 6000] loss: 0.6605622554301299
[SGD, 5, 8000] loss: 0.6688320308341936
[SGD, 5, 10000] loss: 0.6604129620817548
[SGD, 5, 12000] loss: 0.640438791558001
Accuracy of the network with SGD optimizer: 75.79%

```

Рис. 5 — Результат работы нейронной сети VGG16 с оптимизатором SGD




Рис. 6 — Результат работы нейронной сети VGG16 с оптимизатором AdaDelta

4 Выводы

В результате выполнения лабораторной работы на PyTorch были реализованы нейронные сети LeNet, VGG16, ResNet (34) с оптимизаторами SGD, AdaDelta, NAG, Adam.

```

[AdaDelta, 1, 2000] loss: 2.2499634036421776
[AdaDelta, 1, 4000] loss: 1.9987514218091964
[AdaDelta, 1, 6000] loss: 1.7931090692132712
[AdaDelta, 1, 8000] loss: 1.6870725436359644
[AdaDelta, 1, 10000] loss: 1.6093845569044352
[AdaDelta, 1, 12000] loss: 1.5544235248938203
[AdaDelta, 2, 2000] loss: 1.4484482859969139
[AdaDelta, 2, 4000] loss: 1.3755523818135262
[AdaDelta, 2, 6000] loss: 1.3382227805368603
[AdaDelta, 2, 8000] loss: 1.2961661068052053
[AdaDelta, 2, 10000] loss: 1.2599063619673252
[AdaDelta, 2, 12000] loss: 1.2402095287442207
[AdaDelta, 3, 2000] loss: 1.149511059306562
[AdaDelta, 3, 4000] loss: 1.0996184348287061
[AdaDelta, 3, 6000] loss: 1.1147168178567664
[AdaDelta, 3, 8000] loss: 1.0707872729040682
[AdaDelta, 3, 10000] loss: 1.0386965093459002
[AdaDelta, 3, 12000] loss: 1.059362464456819
[AdaDelta, 4, 2000] loss: 0.9536775740922895
[AdaDelta, 4, 4000] loss: 0.9114514478055061
[AdaDelta, 4, 6000] loss: 0.9339479059856385
[AdaDelta, 4, 8000] loss: 0.9156570988511666
[AdaDelta, 4, 10000] loss: 0.8876173911169172
[AdaDelta, 4, 12000] loss: 0.9052881841888738
[AdaDelta, 5, 2000] loss: 0.7754616069906916
[AdaDelta, 5, 4000] loss: 0.7705171618932727
[AdaDelta, 5, 6000] loss: 0.7868890565221663
[AdaDelta, 5, 8000] loss: 0.7946296058123989
[AdaDelta, 5, 10000] loss: 0.803992246422873
[AdaDelta, 5, 12000] loss: 0.794486996242762
Accuracy of the network with AdaDelta optimizer: 70.08%

```

Рис. 7 — Результат работы нейронной сети VGG16 с оптимизатором AdaDelta




Рис. 8 — Результат работы нейронной сети VGG16 с оптимизатором NAG

```

[NAG, 1, 2000] loss: 2.2052371264100077
[NAG, 1, 4000] loss: 1.8782774092555046
[NAG, 1, 6000] loss: 1.7589381381571292
[NAG, 1, 8000] loss: 1.6578851858526469
[NAG, 1, 10000] loss: 1.5526838758885861
[NAG, 1, 12000] loss: 1.4596803367435933
[NAG, 2, 2000] loss: 1.3559429426640273
[NAG, 2, 4000] loss: 1.2919829160198568
[NAG, 2, 6000] loss: 1.218425358850509
[NAG, 2, 8000] loss: 1.1682913989657535
[NAG, 2, 10000] loss: 1.1124468051381409
[NAG, 2, 12000] loss: 1.0649391893462743
[NAG, 3, 2000] loss: 0.9620053204838187
[NAG, 3, 4000] loss: 0.9435190774411895
[NAG, 3, 6000] loss: 0.9309508016603067
[NAG, 3, 8000] loss: 0.892619700523559
[NAG, 3, 10000] loss: 0.8931491761100478
[NAG, 3, 12000] loss: 0.8834547373376553
[NAG, 4, 2000] loss: 0.7568593270430575
[NAG, 4, 4000] loss: 0.7319810788037721
[NAG, 4, 6000] loss: 0.7572539645583892
[NAG, 4, 8000] loss: 0.7354412207430577
[NAG, 4, 10000] loss: 0.728155995545967
[NAG, 4, 12000] loss: 0.7286825955078821
[NAG, 5, 2000] loss: 0.6169509365955987
[NAG, 5, 4000] loss: 0.6194887757776887
[NAG, 5, 6000] loss: 0.610209523483878
[NAG, 5, 8000] loss: 0.6251662377947504
[NAG, 5, 10000] loss: 0.6050609199445899
[NAG, 5, 12000] loss: 0.6192143925252167
Accuracy of the network with NAG optimizer: 77.76%

```

Рис. 9 — Результат работы нейронной сети VGG16 с оптимизатором NAG




Рис. 10 — Результат работы нейронной сети VGG16 с оптимизатором Adam

```

[Adam, 1, 2000] loss: 2.062741495132446
[Adam, 1, 4000] loss: 1.7933170551359654
[Adam, 1, 6000] loss: 1.6396787562668325
[Adam, 1, 8000] loss: 1.5327414567321538
[Adam, 1, 10000] loss: 1.4783946781754493
[Adam, 1, 12000] loss: 1.4157332792580128
[Adam, 2, 2000] loss: 1.3190140459835529
[Adam, 2, 4000] loss: 1.275151315279305
[Adam, 2, 6000] loss: 1.2704246242195367
[Adam, 2, 8000] loss: 1.2003670376949012
[Adam, 2, 10000] loss: 1.1730571854375302
[Adam, 2, 12000] loss: 1.1478230611942708
[Adam, 3, 2000] loss: 1.0443953196927904
[Adam, 3, 4000] loss: 1.0508131722975522
[Adam, 3, 6000] loss: 1.0266857065763324
[Adam, 3, 8000] loss: 1.002074948183261
[Adam, 3, 10000] loss: 0.9888059260174632
[Adam, 3, 12000] loss: 0.9903911812938749
[Adam, 4, 2000] loss: 0.8600541572719812
[Adam, 4, 4000] loss: 0.8675965261529199
[Adam, 4, 6000] loss: 0.8381941680572927
[Adam, 4, 8000] loss: 0.8556119765313343
[Adam, 4, 10000] loss: 0.8459248103424907
[Adam, 4, 12000] loss: 0.8205485811084509
[Adam, 5, 2000] loss: 0.6895021487821359
[Adam, 5, 4000] loss: 0.6896921383730369
[Adam, 5, 6000] loss: 0.6925000499776798
[Adam, 5, 8000] loss: 0.6893276025557424
[Adam, 5, 10000] loss: 0.6943560376940295
[Adam, 5, 12000] loss: 0.6944746703179553
Accuracy of the network with Adam optimizer: 67.48%

```

Рис. 11 — Результат работы нейронной сети VGG16 с оптимизатором Adam

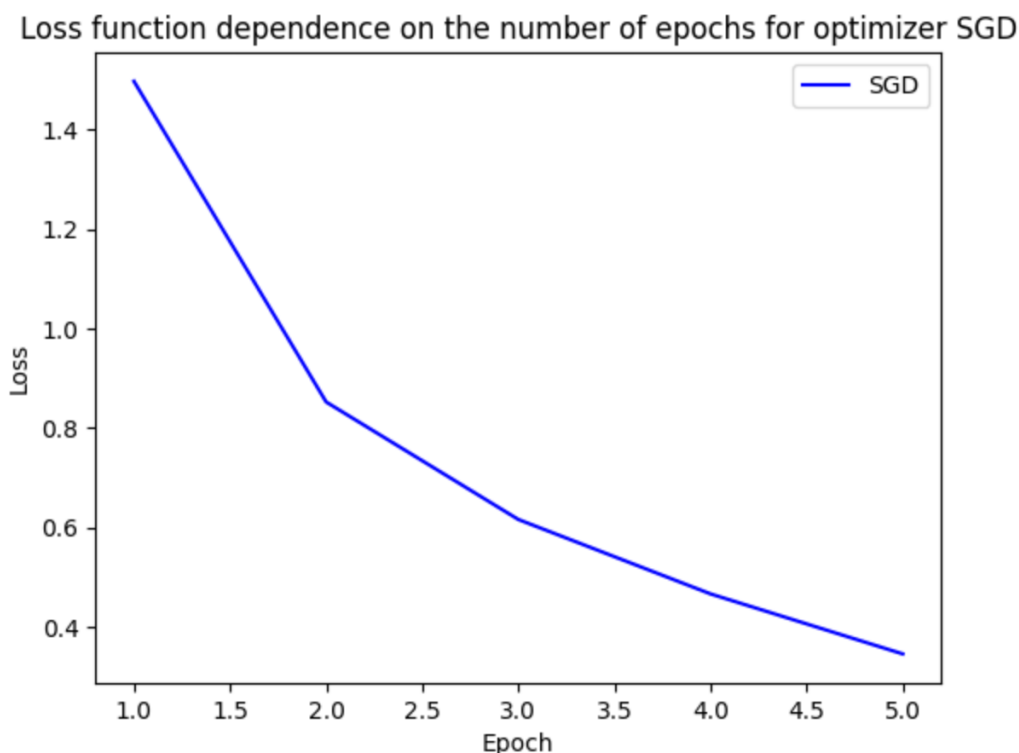


Рис. 12 — Результат работы нейронной сети ResNet (34) с оптимизатором SGD

```

Epoch 1, Optimizer: SGD, Loss: 1.4968267785923561
Epoch 2, Optimizer: SGD, Loss: 0.8519377305227167
Epoch 3, Optimizer: SGD, Loss: 0.6156465031606767
Epoch 4, Optimizer: SGD, Loss: 0.4660088890958625
Epoch 5, Optimizer: SGD, Loss: 0.3453148864114376
Accuracy of the network on the test images with SGD optimizer: 74.23%

```

Рис. 13 — Результат работы нейронной сети ResNet (34) с оптимизатором SGD

Loss function dependence on the number of epochs for optimizer Adadelta




Рис. 14 — Результат работы нейронной сети ResNet (34) с оптимизатором AdaDelta

```
Epoch 1, Optimizer: Adadelta, Loss: 1.2652742084296769
Epoch 2, Optimizer: Adadelta, Loss: 0.671191810532604
Epoch 3, Optimizer: Adadelta, Loss: 0.5188204992343398
Epoch 4, Optimizer: Adadelta, Loss: 0.4174706535723508
Epoch 5, Optimizer: Adadelta, Loss: 0.3345142348438425
Accuracy of the network on the test images with Adadelta optimizer: 68.98%
```

Рис. 15 — Результат работы нейронной сети ResNet (34) с оптимизатором AdaDelta

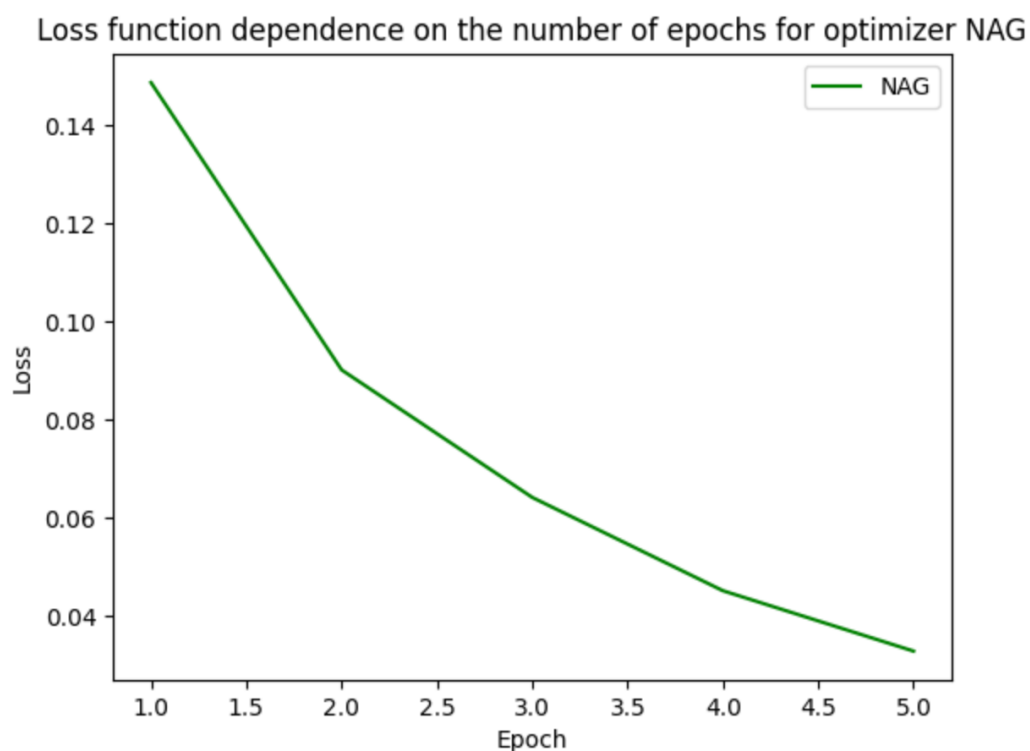


Рис. 16 — Результат работы нейронной сети ResNet (34) с оптимизатором NAG

```
Epoch 1, Optimizer: NAG, Loss: 0.1486827559028383
Epoch 2, Optimizer: NAG, Loss: 0.09016190473314213
Epoch 3, Optimizer: NAG, Loss: 0.06419006415614215
Epoch 4, Optimizer: NAG, Loss: 0.04518884563829888
Epoch 5, Optimizer: NAG, Loss: 0.03288775410623676
Accuracy of the network on the test images with NAG optimizer: 87.33%
```

Рис. 17 — Результат работы нейронной сети ResNet (34) с оптимизатором NAG




Рис. 18 — Результат работы нейронной сети ResNet (34) с оптимизатором Adam


```
Epoch 1, Optimizer: Adam, Loss: 1.529735531495965  
Epoch 2, Optimizer: Adam, Loss: 0.926960490434371  
Epoch 3, Optimizer: Adam, Loss: 0.6846035719298951  
Epoch 4, Optimizer: Adam, Loss: 0.5517108656103958  
Epoch 5, Optimizer: Adam, Loss: 0.45280647239721644  
Accuracy of the network on the test images with Adam optimizer: 80.00%
```

Рис. 19 — Результат работы нейронной сети ResNet (34) с оптимизатором Adam